

# HyCache: a User-Level Caching Middleware for Distributed File Systems

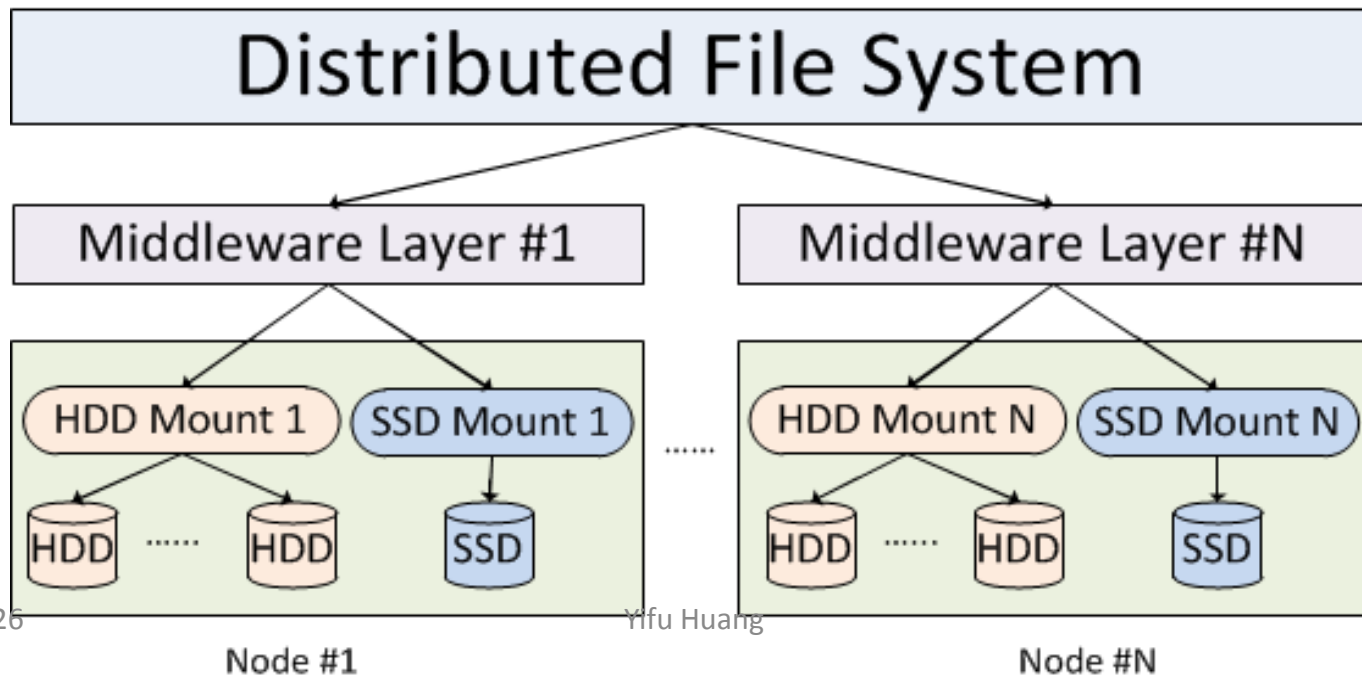
Dongfang Zhao, Ioan Raicu  
IPDPS workshop 2013

# Motivation

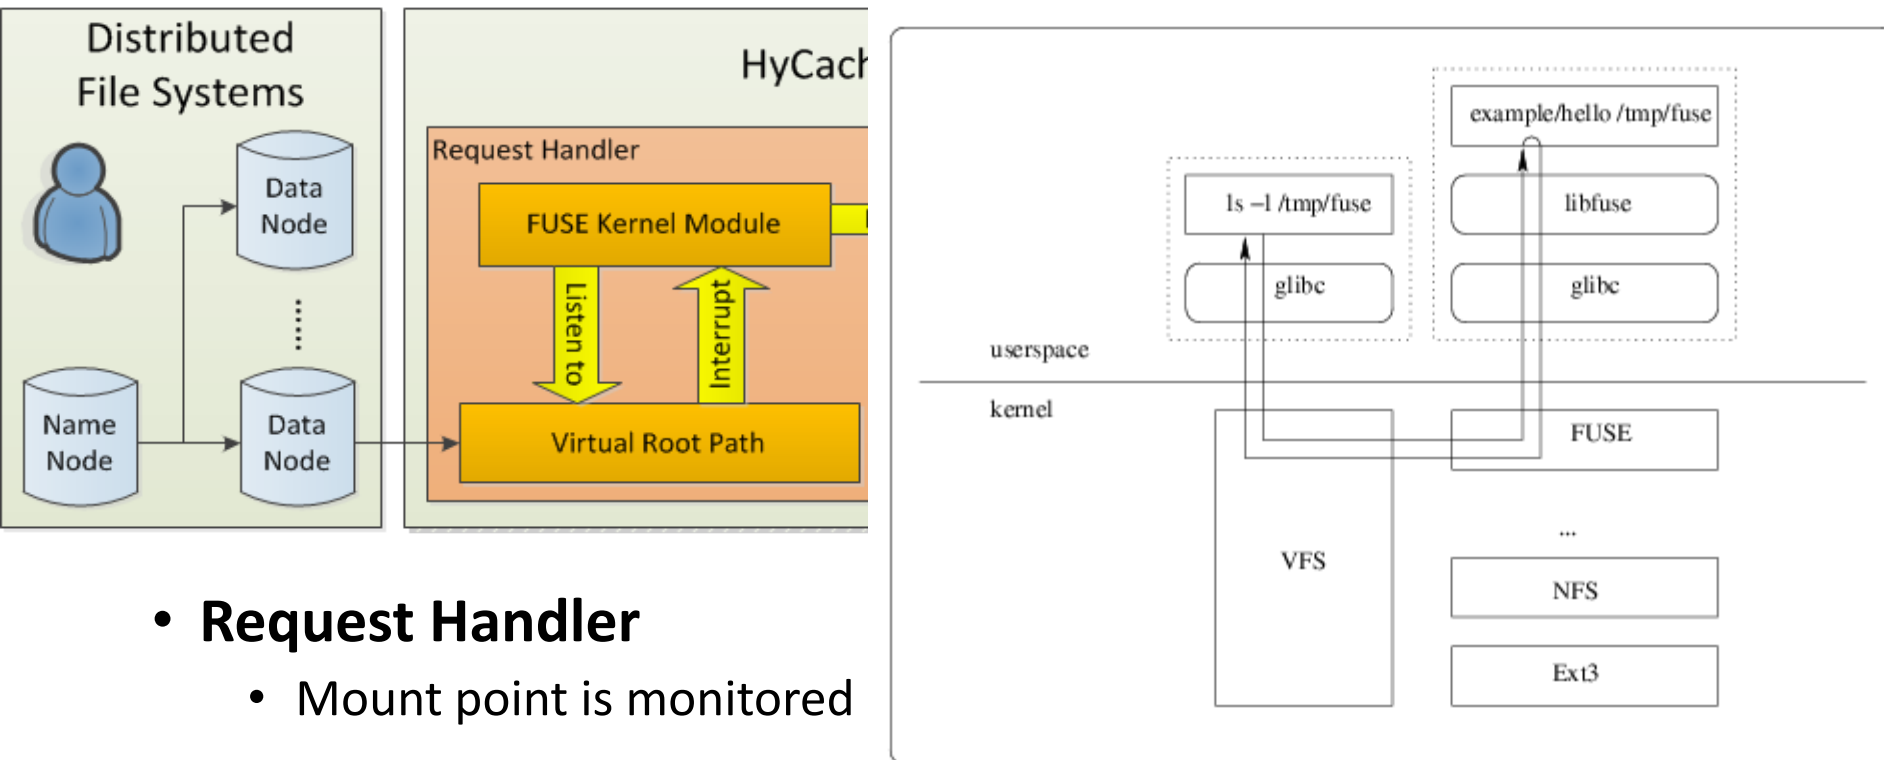
- One of the bottlenecks of distributed file systems is mechanical hard disk drives (**HDD**)
  - **Slow** increase in bandwidth
  - **Slow** decrease in latency
  - **Exponential** increase in capacity
- **Unbalanced!**
- SSD could bridge the gap nicely between RAM and HDD for a distributed file system
- **SSD performance + capacity of HDD**

# Middleware hierarchy

- The storage hierarchy with a middleware between **distributed file systems** and **local file systems**
- Manage **heterogeneous** storage devices for distributed file systems

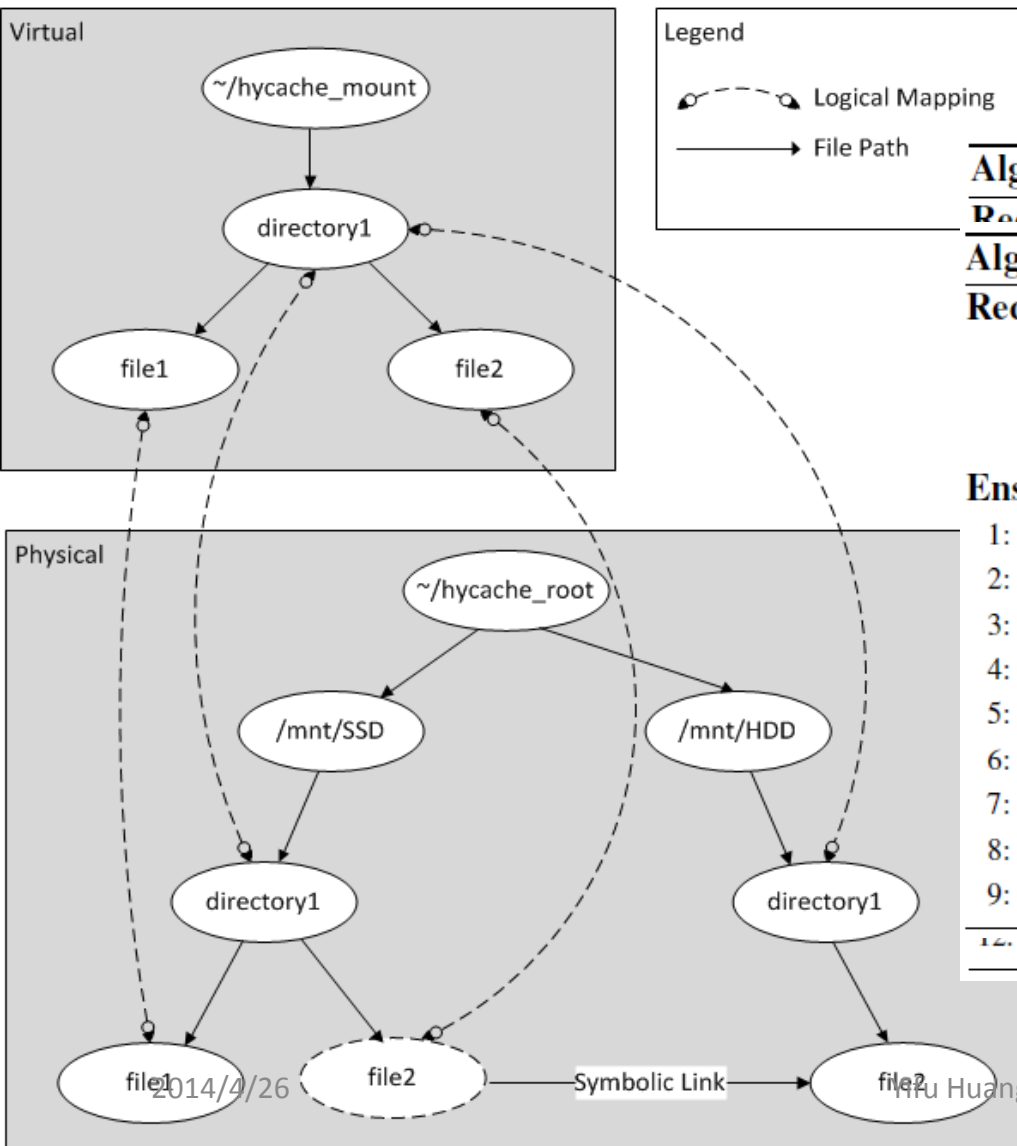


# HyCache architecture



- **Request Handler**
  - Mount point is monitored
- **File Dispatcher**
  - File manipulations and cache algorithms
- **Data Manipulator**
  - Manipulates data between two logical access points

# HyCache implementation



## Algorithm 1 Open a file in HyCache

**Require:**  $F$  is the file requested by the end user;  $Q$  is the

## Algorithm 3 Rename a file in HyCache

**Require:**  $F$  is the file requested by the end user to rename;  
 $F'$  is the new file name;  $Q$  is the queue used for the replacement policy; SSD is the mount point of SSD drive; HDD is the mount point of HDD drive


**Ensure:**  $F$  is renamed to  $F'$

- 1: **if**  $F$  is a symbolic link in SSD **then**
- 2:   rename  $F$  to  $F'$  in HDD
- 3:   remove  $F$  in SSD
- 4:   create the symbolic link  $F'$  in SSD
- 5: **else**
- 6:   rename  $F$  to  $F'$  in SSD
- 7:   rename  $F$  to  $F'$  in  $Q$
- 8: **end if**
- 9: update  $F'$  position in  $Q$

12: **open**  $F$  in SSD

# HyCache implementation (Cont.)

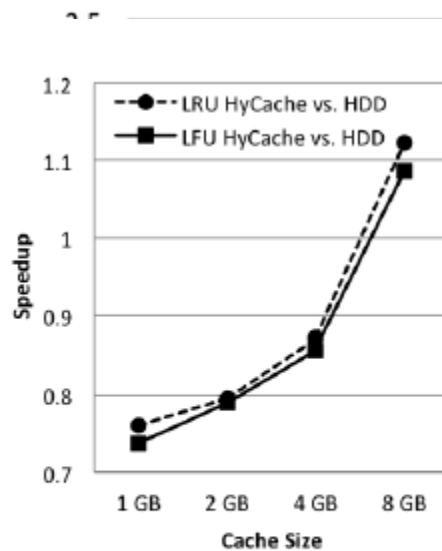
- Two built-in cache algorithms: **LRU** and **LFU**
  - Free to plug in other cache algorithms
  - With the standard C library <search.h>
  - Doubly-linked list
  - Element: filename, access time, number of access, etc
- LRU queue in HyCache



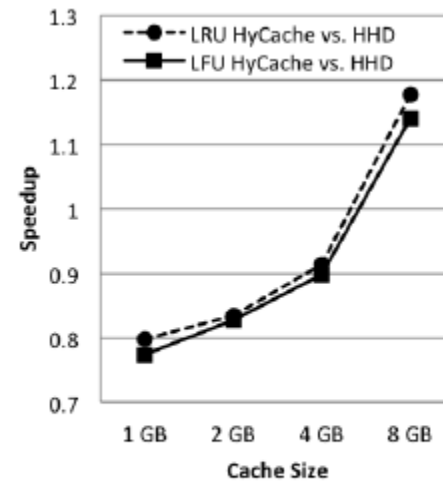
# Evaluation

- Tested the **functionality** and **performance** of HyCache in four experiments
  - First two are benchmarks with synthetic data to test the raw bandwidth of HyCache
    - Micro-benchmarks
    - macro-benchmarks
  - The third and fourth experiments are to test the functionality of HyCache with a real application
    - MySQL
    - HDFS

# Evaluation (Cont.)



(a) HyCache vs. HDD



(b) HyCache vs. HDD



# Contribution

- Designed and implemented HyCache
  - High throughput, low latency, strong consistency, single namespace, and multithreaded support
- Developed a middleware layer
  - Delivered 28% improvement in HDFS performance
- Extensive performance evaluation
  - Competitive with kernel-level file systems

# HyCache+: Towards Scalable High-Performance Caching Middleware for Parallel File Systems

Dongfang Zhao, Kan Qiao, Ioan Raicu  
CCGrid 2014

# Motivation

- The ever-growing **gap** between the **computation** and **I/O** is one of the fundamental challenges for today's large scale computing systems
  - The number of compute cores follows Moore's Law
  - The storage systems have been improved at a much slower pace
- Even larger for modern high-performance computing (**HPC**) systems

# HyCache+

- Distributed storage middleware HyCache+
  - Allows I/O to effectively leverage the **high bi-section bandwidth** of the **high-speed interconnect** of massively parallel high-end computing systems
  - The primary place for holding hot data for the applications
  - Only asynchronously swaps with cold data on the remote parallel file system
- Opens the door to providing both **high performance** and cost-effective **large capacity**

# HyCache+ hierarchy




Table I

SOME KEY HYCACHE+ IMPROVEMENTS OVER HYCACHE

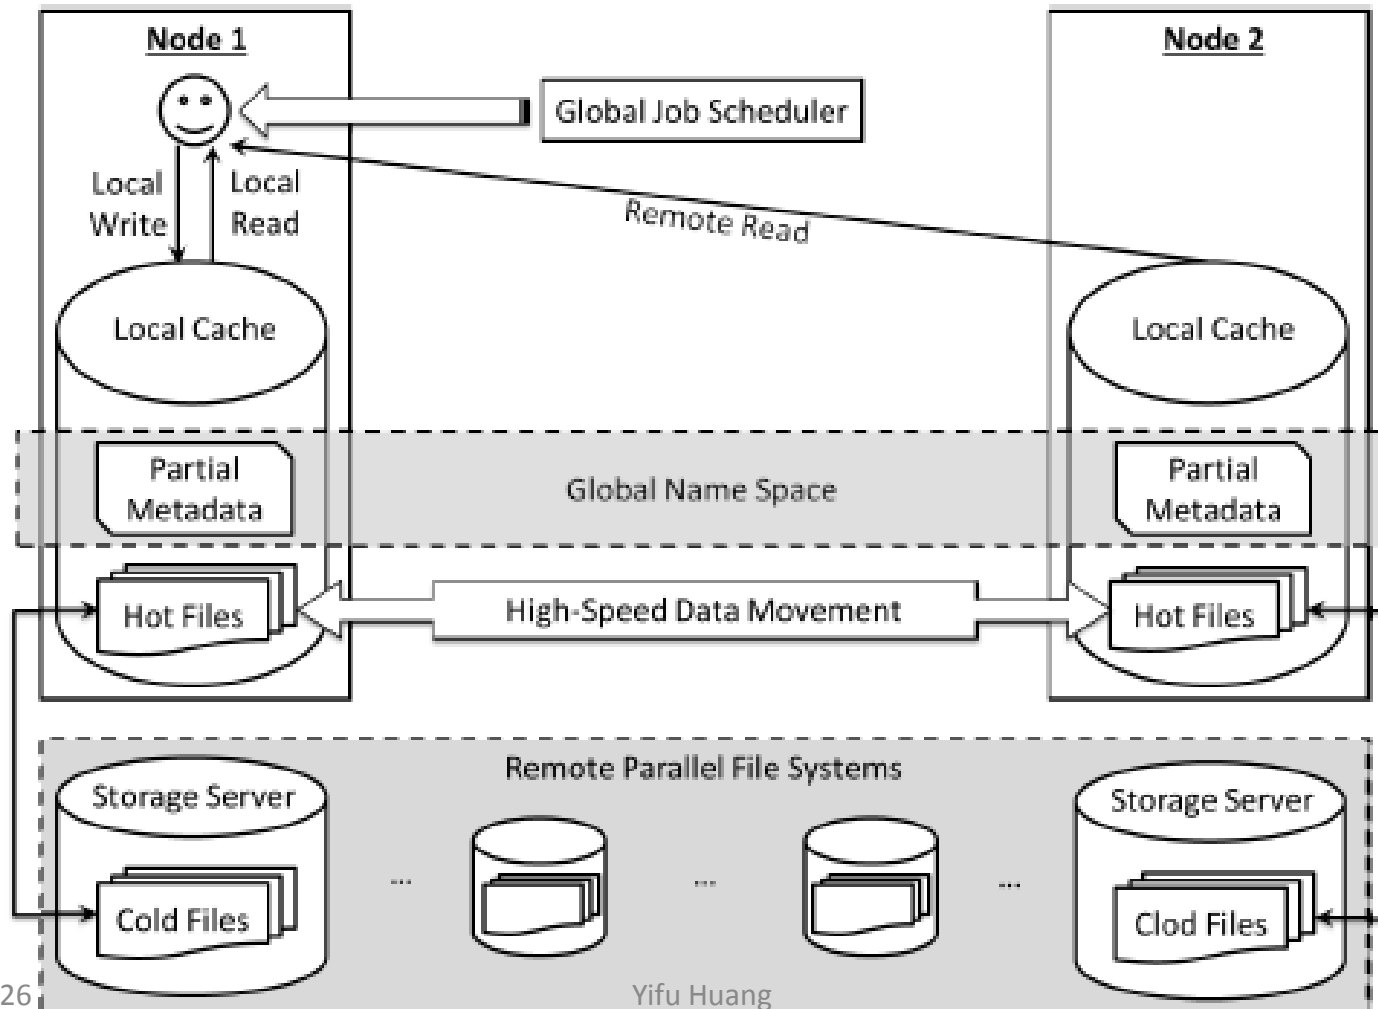
Mechanism	HyCache+	HyCache
Network Storage	Yes	No
Data Movement	Local & Remote	Local Only
Replica	Arbitrary (e.g. 3)	1
Scalability	4096-cores	1-node
Metadata	DHT	Symbolic Link

□

- 1 – Cache hit
- 2 – Data swap between cache and GPFS
- 3 – Data movement across compute nodes



# HyCache+ design overview



# HyCache+ Metadata

- **DHT** is the translator between local partial metadata and the global namespace

Key	Value
~/	drwxrwxr-x; 4.0K; ~/homedir/subdir
~/homedir/	drwxrwxr-x; 4.0K; ~/homedir/subdir, ~/homedir/homefile
~/homedir/subdir/	drwxrwxr-x; 4.0K; ~/homedir/subdir/subfile
~/homedir/homefile	-rw-rw-r--; 423M; Node 1
~/homedir/subdir/subfile	-rw-rw-r--; 133M; Node 2
⋮	⋮

- Only the primary copy is used for read and write operation, and replicas are only used to avoid data loss in case of node failures

# HyCache+ Fault Tolerance

- Synchronous replicas
  - A costly method, and satisfies the strong consistency requirement, if needed
- Asynchronous replicas
  - Would deliver the highest throughput, while being compromised on the possibility of failing to recover before the asynchronous update is completed
- Erasure coding
  - Trades off between performance and consistency



# 2-Layer Scheduling

That is to solve the objective function

$$\arg \min_Q \sum_{A_k \in A} \sum_{M_l \in M} \sum_{F_i \in F^k} \sum_{M_j \in M} \text{Size}(F_i) \cdot P_{i,j} \cdot Q_{k,l},$$

## • 1. Job Scheduling

---

**Algorithm 1** Global Schedule

---

**Input:** The  $x^{th}$  job to be scheduled

**Output:** The  $y^{th}$  machine where the  $x^{th}$  job should be scheduled

```

1: function GLOBALSCHEDULE( $x$ )
2:    $MinCost \leftarrow \infty$ 
3:    $y \leftarrow \text{NULL}$ 
4:   for  $M_i \in M$  do
5:      $Cost \leftarrow 0$ 
6:     for  $F_j \in F^x$  do
7:       Find  $M_k$  such that  $P_{j,k} = 1$ 
8:        $Cost \leftarrow Cost + \text{Size}(F_j)$ 
9:     end for
10:    if  $Cost < MinCost$  then
11:       $MinCost \leftarrow Cost$ 
12:       $y \leftarrow i$ 
13:    end if
14:  end for
15:  return  $y$ 
16: end function

```

$$\sum_{M_j \in M} P_{i,j} = 1, \forall F_i \in F,$$

$$\sum_{M_j \in M} Q_{i,j} = 1, \forall A_i \in A,$$

$$P_{i,j}, Q_{i,j} \in \{0, 1\}, \forall i, j.$$

by $A_k \in A$
n $M_j \in M$
d on $M_j \in M$

# 2-Layer Scheduling (Cont.)

- **2. Heuristic Caching**

- The problem of finding optimal caching on multiple-disk is proved to be NP-hard
- Propose a heuristic algorithm of  **$O(n \lg n)$**
- **Cost:** the to-be-evicted file size multiplied by its access frequency after the current processing position in the reference sequence
- **Gain:** the to-be-cached file size multiplied by its access frequency after the current fetch position in the reference sequence

# 2-Layer Scheduling (Cont.)

- **2. Heuristic Caching (Cont.)**

- **Rule 1.** Every fetch should bring into the cache the very next file in the reference sequence if it is not yet in the cache.
- **Rule 2.** Never fetch a file to the cache if the total **cost** of the to-be-evicted files is greater than the **gain** of fetching this file.
- **Rule 3.** Every fetch should discard the files in the increasing order of their **cost** until there is enough space for the newly fetched file. If the cache has enough space for the new file, no eviction is needed.

# 2-Layer Scheduling (Cont.)

## • 2. Heuristic Caching (Cont.)

### • Example

$R = (r_1, r_2, r_3, r_4, r_5, r_6, r_7, r_8, r_9)$ . Let  $File(r_1) = F_1$ ,  $File(r_2) = F_2$ ,  $File(r_3) = F_3$ ,  $File(r_4) = F_4$ ,  $File(r_5) = F_3$ ,  $File(r_6) = F_1$ ,  $File(r_7) = F_2$ ,  $File(r_8) = F_4$ ,  $File(r_9) = F_3$ , and  $Size(F_1) = 20$ ,  $Size(F_2) = 40$ ,  $Size(F_3) = 9$ ,  $Size(F_4) = 40$ . Let the cache capacity be 100. According to *Rule 1*, the first three files to be fetched to cache are  $(F_1, F_2, F_3)$ . Then we need to decide if we want to fetch  $F_4$ . Let  $Cost(F_i)$  be the cost of evicting  $F_i$ . Then we have  $Cost(F_1) = 20 \times 1 = 20$ ,  $Cost(F_2) = 40 \times 1 = 40$ , and  $Cost(F_3) = 9 \times 2 = 18$ . According to *Rule 3*, we sort the costs in the increasing order  $(F_3, F_1, F_2)$ . Then we evict the files in the sorted list, until there is enough room for the newly fetched file  $F_4$  of size 40. In this case, we only need to evict  $F_3$ , so that the free cache space is  $100 - 20 - 40 = 40$ , just big enough for  $F_4$ . Before replacing  $F_3$  by  $F_4$ , *Rule 2* is referred to ensure that the cost is smaller than the gain, which is true in this case by observing that the gain of prefetching  $F_4$  is 40, larger than  $Cost(F_3) = 18$ .

# 2-Layer Scheduling (Cont.)

- **2. Heuristic Caching (Cont.)**

- Fetch a file to cache or processor

---

**Algorithm 2** Fetch a file to cache or processor

---

**Input:**  $i$  is the reference index being processed

```
1: procedure FETCH( $i$ )
2:   if  $\{r_j | File(r_j) = File(r_{i+1}) \wedge j > i + 1\} \neq \emptyset$  then
3:      $flag, D \leftarrow GetFilesToDiscard(i, i + 1)$ 
4:     if  $flag = successful$  then
5:       Evict  $D$  out of the cache
6:       Fetch  $File(r_{i+1})$  to the cache
7:     end if
8:   end if
9:   Access  $File(r_{i+1})$  (either from the cache or the disk)
10: end procedure
```

---

# 2-Layer

## • 2. Heuristics

- Get set

---

### Algorithm 3 Get set of files to be discarded

---

**Input:**  $i$  is the reference index being processed;  $j$  is the reference index to be (possibly) fetched to cache

**Output:** *successful* –  $File(r_j)$  will be fetched to the cache and  $D$  will be evicted; *failed* –  $File(r_j)$  will not be fetched to the cache

```


1: function GETFILESTODISCARD( $i, j$ )
2:   if  $Size(S) + Size(File(r_j)) \leq C$  then
3:     return successful,  $\emptyset$ 
4:   end if
5:    $num \leftarrow$  Number of occurrences of  $File(r_j)$  from  $j+1$ 
6:    $gain \leftarrow num \cdot Size(File(r_j))$ 
7:    $cost \leftarrow 0$ 
8:    $D \leftarrow \emptyset$ 
9:   Sort the files in  $S$  in the increasing order of the cost
10:  for  $F \in S$  do
11:     $tot \leftarrow$  Number of references of  $F$  from  $i+1$ 
12:     $cost \leftarrow cost + tot \cdot Size(F)$ 
13:    if  $cost < gain$  then
14:       $D \leftarrow D \cup \{F\}$ 
15:    else
16:       $D \leftarrow \emptyset$ 
17:      return failed,  $D$ 
18:    end if
19:    if  $Size(S \setminus D) + Size(File(r_j)) \leq C$  then
20:      break
21:    end if
22:  end for
23:  return successful,  $D$ 
24: end function

```

# Evaluation

- 1024 nodes (4096 cores)

• E<sub>2</sub>



nd

# Contribution

- Design and implement a scalable high-performance caching middleware, namely HyCache+
  - Improve the I/O performance
- Propose and analyze a novel caching approach —2-Layer Scheduling (2LS)
  - Optimize the network cost
  - Heuristically reduce the disk I/O cost
- Evaluate at large scale
  - Report their performance on a leadership class supercomputer